

Actuation-Aware Sensor Validation for Fault-Resilient Closed-Loop Hydroponic Control

MD. Sameer Iqbal Chowdhury^a, Mikdam-Al-Maad Ronoue^a, Md Asaduzzaman^b, Lafifa Jamal^{a,*}

^a*Department of Robotics and Mechatronics Engineering, University of Dhaka, Dhaka-1000, Bangladesh*

^b*Department of Agriculture and Food Technology, Kyoto University of Advanced Sciences, Japan*

Abstract

This paper proposes AASVR-R, a streaming measurement-to-actuation reliability layer for closed-loop hydroponic control. Unlike a conventional detector, AASVR-R can record a suspect measurement while withholding its authority to drive an actuator. The novelty is the auditable supervisor interface: validation, trusted-state update, rectification, authorization, and response verification are separate decisions, while standard robust gates and Page/CUSUM residual tests provide the signal-domain evidence. Held-out replay on a real NFT hydroponic deployment shows competitive detection against tuned filtering, process-monitoring, model-based, and one-class baselines while preserving low false-authorized actuation and a traceable decision log. Counterfactual response and diagnostic-excitation replays are treated as capability demonstrations, not measured closed-loop evidence. The results identify actuator-state logging and reference sensing as the principal requirements for resolving plausible stuck values in future deployments.

Keywords: sensor validation, data rectification, fault detection, process control, hydroponics, actuation

1. Introduction

Hydroponic controlled-environment systems depend on continuous measurement and actuation. In a recirculating Nutrient Film Technique (NFT)

*Corresponding author.

installation, the streams that describe pH, electrical conductivity (EC), air temperature, humidity, water temperature, water level, and carbon dioxide (CO₂) can also trigger dosing pumps, chillers, valves, cooling devices, and circulation equipment. A faulty measurement is therefore not only a data-quality issue; it can become a physical command that perturbs the nutrient solution or growth environment.

Standard filtering and anomaly-detection methods usually score the signal. Closed-loop operation also requires a second question: whether the current value is trustworthy enough to update the controller state or authorize an actuator. A single high-pH spike should not authorize acid dosing, and a temporary EC excursion should not be treated the same as a persistent concentration shift. This motivates actuation-aware sensor validation, where logging, state update, rectification, and actuation authorization are separate decisions.

We present a closed-loop NFT hydroponic platform and propose AASVR-R, an actuation- and response-aware sensor-validation layer, as its fault-resilience mechanism. The contribution is structural: AASVR-R formalizes a supervisory measurement-to-actuation interface in which robust gates decide whether a sample is plausible, a state machine decides whether it can update the trusted state, a rectifier decides what value is exposed to control, an authorization rule decides whether actuation is allowed, and a response residual checks whether the measured process reacts as expected after an authorized command. The contributions are:

1. an integrated hydroponic deployment dataset and sensor-to-actuator formulation for NFT control;
2. AASVR-R, a streaming supervisory interface that couples standard validation evidence, bounded rectification, risk-aware authorization, and actuator-response reliability memory;
3. a held-out synthetic-fault protocol with validation-only tuning, tuned baselines, ablations, component analysis, and paired statistical tests;
4. a reproducible Docker and LaTeX workflow that regenerates the tables, figures, and diagnostic outputs.

2. Related Work

2.1. Sensor faults in closed-loop systems

Sensor data-quality reviews distinguish missingness, outliers, drift, stuck values, saturation, and inconsistent rates as recurring faults in deployed sys-

tems [1]. Agricultural IoT systems face the same problem when sensor failure or miscalibration propagates into monitoring or control decisions [2, 3]. In closed-loop hydroponics, these faults have downstream consequences because an erroneous measurement can change the actuator command. This motivates treating validation, rectification, and actuation as separate decisions rather than as one smoothing step.

2.2. Filtering and monitoring baselines

The baseline family used in this paper covers common online filters, process-monitoring methods, model-based statistical tests, and lightweight one-class anomaly detectors. Hampel-style robust outlier detection uses robust scale estimates to suppress isolated deviations [4]. Kalman filtering gives a probabilistic recursive estimator for linear dynamical systems [5]. EWMA and CUSUM-style monitoring are standard approaches for persistent shifts [6, 7]. Generalized likelihood ratio (GLR) tests provide a classical statistical framework for abrupt change detection [8]. Analytical redundancy, parity/observer residuals, and reconfiguration logic are central themes in fault-detection and isolation (FDI) for control systems [9, 10, 11]; AASVR-R does not replace these model-based methods, but supplies the controller-facing supervisor that decides whether a measurement or residual may drive actuation. PCA and recursive PCA-style residual monitoring use residual statistics to capture coordinated changes or adaptive operating regimes [12, 13, 14]. Isolation Forest, One-Class SVM, and local outlier factor add common one-class machine-learning comparisons for anomaly detection [15, 16, 17]. Existing hydroponic automation studies have demonstrated reliable pH/EC regulation and IoT monitoring, but they generally treat validation as part of measurement filtering or rule execution rather than as a separate authorization decision [18, 19]. AASVR-R is complementary to these methods because it adds a control-facing authorization interface.

2.3. External process datasets

Public process datasets are useful stress tests for measurement validation, but their labels do not necessarily represent hydroponic sensor faults. HAI 21.03 contains industrial-control attack data from a hardware-in-the-loop testbed [20], and SKAB contains multivariate process-loop experiments with anomaly and changepoint labels [21]. These datasets are retained in the reproducibility artifacts as supplementary transfer checks, not as headline evidence. SWaT, WaDi, and Tennessee Eastman remain relevant larger

benchmarks for future validation under their respective access and redistribution constraints [22, 23, 24].

3. Hydroponic System Design and Deployment

The target platform is a closed NFT hydroponic system for lettuce growth. NFT systems recirculate a shallow nutrient stream through sloped growth channels, allowing water and nutrient reuse when solution management is reliable [25, 26]. The deployed system is organized into two physical modules. The Growth Module houses the plants, lighting, air-temperature and humidity sensing, CO₂ sensing, visual monitoring, and environmental actuation. The Nutrient Solution Container Module stores and recirculates the nutrient solution and contains the water-quality sensors and nutrient-solution actuators.




Figure 1: Hydroponic system architecture and AASVR-R location in the sensing-to-actuation path.

The Growth Module uses a DHT22 sensor for air temperature and humidity and an MH-Z19C nondispersive infrared CO₂ sensor for gas concentration. CO₂ regulation is implemented through a normally closed pneumatic solenoid valve connected to a CO₂ cylinder. Air-temperature regulation is supported by a thermoelectric cooler. Two 40 W LED grow-light arrays provide a controlled light cycle, and a camera provides visual monitoring of the crop during the growth cycle.

The Nutrient Solution Container Module uses an Atlas Scientific analog pH kit, an Atlas Scientific Mini Conductivity K 1.0 kit for EC, a water-temperature sensor, and a JSN-SR04T waterproof ultrasonic distance sensor



Figure 2: Physical hydroponic system during operation.

for water level. pH correction is implemented through acid/base dosing, and EC correction through concentrated nutrient solution or distilled-water addition. After a pH or EC dosing action, the dosing subsystem is disabled for a 10-minute mixing interval before further dosing is allowed. A refrigerative water chiller regulates nutrient-solution temperature, an aerator runs throughout the growth cycle to support dissolved oxygen, and a water pump recirculates nutrient solution through the NFT channels.

The control unit connects sensors to an Arduino Nano and actuators to an Arduino Mega. Both microcontrollers communicate with a Raspberry Pi 4 Model B, which serves as the central control computer. This sensor-microcontroller-edge-computer structure is consistent with prior hydroponic automation systems, but the present paper focuses on the measurement-to-actuation reliability layer rather than on dashboarding as a contribution [18, 19]. The Raspberry Pi receives median-filtered sensor readings from the Arduino Nano, logs the readings, updates the ThingSpeak-backed dashboard, applies the validation/control logic, and sends authorized actuation commands to the Arduino Mega. The dashboard and camera feed are retained as deployment infrastructure; they are not claimed as the research contribution of this paper.

Table 1: Hydroponic platform components relevant to measurement validation and actuation. Water level is listed because it is part of the deployed system, but it is excluded from the primary quantitative analysis because the available raw values are not defensibly calibrated.

Subsystem	Measurements	Actuation or use
Growth module	Air temperature, humidity	Thermoelectric cooling
Growth module	CO2 concentration	CO2 solenoid valve
Growth module	Camera feed	Visual crop monitoring
Growth module	Light schedule	LED grow lights
Solution module	pH	Acid/base dosing, 10 min lockout
Solution module	EC	Nutrient/water dosing, 10 min lockout
Solution module	Water temperature	Refrigerative chiller
Solution module	Water level	Warning only in this analysis
Solution module	Recirculation state	Pump and aerator operation

The evaluation below treats this system as a deployed measurement-to-actuation stream: raw measurements are replayed through validation methods, transformed into trusted states, and then passed through an authorization rule before any actuator command is counted.

4. Problem Formulation

For hydroponic sensor s at time t , let $y_s(t)$ be the raw measurement and $\hat{x}_s(t)$ be the trusted value exposed to control. The method observes recent history $Y_s(t - w : t)$, broad physical limits $[P_{s,\min}, P_{s,\max}]$, a control band $[L_s^a, U_s^a]$, a rate limit r_s , and any available actuator state $a(t)$. In this system, s may correspond to pH, EC, humidity, air temperature, water temperature, or CO2, while $a(t)$ may represent dosing, cooling, circulation, or gas-release activity.

Table 2: Notation used in the AASVR-R formulation. Variables are defined for one sensor stream and then applied independently or jointly across the hydroponic sensor set.

Symbol	Meaning
$y_s(t)$	Raw value from sensor s at sample t
$\hat{x}_s(t)$	Trusted value exposed to the controller
$Y_s(t - w : t)$	Rolling window of recent raw values
$[P_{s,\min}, P_{s,\max}]$	Broad physically plausible range
$[L_s^a, U_s^a]$	Actuation or control band
r_s	Maximum plausible rate of change
u_s	Sensor uncertainty or minimum measurement resolution term
$q_s(t)$	Trust score in $[0, 1]$
$z_s(t)$	Sensor state in the AASVR-R state set
$h_s(t)$	Evidence from the recent persistence/confirmation window
$A_s(t)$	Binary actuation authorization flag
θ	Tunable AASVR-R parameter vector

We define a measurement as *untrusted* when it is missing, physically implausible, rate-inconsistent, inconsistent with recent trusted state, or inconsistent with expected actuator effects. We define a *false actuation* as an authorization caused by untrusted evidence. A *missed actuation* is a failure to authorize action when the trusted state remains outside the control band.

The validation decision is separated from the control decision. Let $g_s(t) \in \{0, 1\}$ denote whether a raw sample passes all measurement gates, and let $z_s(t)$ denote one of five states: normal, suspect transient, persistent deviation, accepted regime shift, and fault alert. The controller receives $\hat{x}_s(t)$ and $A_s(t)$, not the raw value alone. This separation is necessary because a value can be physically possible but still too uncertain to authorize actuation.

The conceptual evaluation objective is

$$J(\theta) = w_1 E_{\text{pass}} + w_2 E_{\text{reject}} + w_3 A_{\text{false}} + w_4 A_{\text{missed}} + w_5 D_{\text{delay}} + w_6 T_{\text{unsafe}} + w_7 C_{\text{compute}}, \quad (1)$$

where the terms respectively measure anomalous pass-through, false rejection, false actuation, missed actuation, detection delay, residence outside safe bands, and computational cost. The validation sweep uses the operative scalar

$$J_{\text{val}}(\theta) = \text{BA}(\theta) - 0.01 A_{\text{false}}(\theta) - 0.0005 F_{\text{alarm}}(\theta) - 0.001 D_{\text{delay}}(\theta), \quad (2)$$

where BA is balanced accuracy, A_{false} is the mean false-authorized actuation count, F_{alarm} is the mean false-alarm event count, and D_{delay} is mean detection delay in samples. Missed actuation, unsafe-band residence, and runtime are logged diagnostics rather than validation-selection terms in this study.

5. Proposed Method: AASVR-R

5.1. Overview

Figure 3 summarizes the AASVR-R pipeline. The method separates five decisions that are often conflated: measurement validation, trusted-state update, rectification, actuation authorization, and response verification. This separation is the core design choice. In the hydroponic system, this means that a suspicious pH value can be stored with a diagnostic flag while acid/base dosing remains suppressed until the value is persistent and trusted, and repeated dosing authority can be removed if the pH stream does not respond after a command.

5.2. Robust plausibility gate

For each sensor, AASVR-R estimates a robust scale over adjacent changes,

$$\sigma_{\Delta,s}(t) = 1.4826 \text{MAD}(\Delta y_s(t - w : t)). \quad (3)$$

The accepted deviation tolerance is

$$\xi_s(t) = \max(\xi_{s,\min}, c_s \sigma_{\Delta,s}(t), r_s \Delta t, u_s). \quad (4)$$




Figure 3: AASVR-R architecture. Measurement validation, trusted-state update, rectification, actuation authorization, and response verification are separate decisions.

A reading must satisfy broad physical range, maximum rate, and trusted-state deviation gates. The three elementary gates are

$$\begin{aligned} G_s^P(t) &= \mathbf{1}\{P_{s,\min} \leq y_s(t) \leq P_{s,\max}\}, \\ G_s^R(t) &= \mathbf{1}\{|y_s(t) - y_s(t-1)|/\Delta t \leq r_s\}, \\ G_s^D(t) &= \mathbf{1}\{|y_s(t) - \hat{x}_s(t-1)| \leq \xi_s(t)\}. \end{aligned} \quad (5)$$

A separate trend guard $G_s^T(t)$ rejects slow monotonic movement when no relevant actuator activity explains the direction of change. A stuck-value guard $G_s^S(t)$ rejects sustained exact repetition after a preceding change, which targets frozen plausible readings without treating ordinary low-variance noise as a fault. To make the drift component less dependent on a single-window monotonicity rule, AASVR-R also uses a two-sided CUSUM residual test. For residual $e_s(t) = y_s(t) - \hat{x}_s(t-1)$,

$$\begin{aligned} C_s^+(t) &= \max\{0, C_s^+(t-1) + e_s(t) - \kappa_s(t)\}, \\ C_s^-(t) &= \max\{0, C_s^-(t-1) - e_s(t) - \kappa_s(t)\}, \end{aligned} \quad (6)$$

where $\kappa_s(t) = \alpha_c \xi_s(t)$ is the allowed drift allowance. The sequential residual gate is

$$G_s^C(t) = \mathbf{1}\{\max(C_s^+(t), C_s^-(t)) \leq h_c \xi_s(t)\}. \quad (7)$$

This is the Page/CUSUM change-detection principle applied to the trusted-state residual: isolated residuals decay, while persistent one-sided departures

accumulate until they are no longer trusted for state update. The validation-selected hydroponic configuration uses $\alpha_c = 0.10$ and $h_c = 0.60$ for the primary sensors. The overall gate is therefore

$$g_s(t) = G_s^P(t)G_s^R(t)G_s^D(t)G_s^T(t)G_s^S(t)G_s^C(t)G_s^M(t), \quad (8)$$

where $G_s^M(t)$ is one only when the value is present and parseable. The broad physical range is deliberately different from the crop-control band: for example, a pH value can be undesirable for lettuce while still being physically possible and therefore not automatically a sensor fault.

5.3. Trust score

AASVR-R also computes a continuous trust score so that control can distinguish weakly plausible values from strongly confirmed values. The implementation uses

$$q_s(t) = 0.25q_s^P(t) + 0.20q_s^R(t) + 0.20q_s^H(t) + 0.20q_s^A(t) + 0.15q_s^M(t), \quad (9)$$

where q_s^P is the physical-range score, q_s^R is the rate score, $q_s^H = \max(0, 1 - 0.2n_s^{\text{fail}})$ is the persistence score, q_s^A is actuator-consistency evidence, and q_s^M is the missingness score. Each component lies in $[0, 1]$. Trusted-deviation, uncommanded-trend, stuck-value, CUSUM-residual, and actuator-response-residual failures cap the score through the gate outcome rather than adding another free weight. The physical pH/EC dosing code also contains a 10-minute mixing lockout. Numerical settings are separated from the method definition and selected through the validation protocol in Section 6.4.

5.4. State-machine persistence logic

Figure 4 shows the five persistence states. A single failed gate moves a sensor into a suspect transient state. Repeated failures create a persistent-deviation state. Consistent evidence can be accepted as a new process regime, while missing, stuck, saturated, or contradictory readings escalate to fault alert. This matters for hydroponics because real mixing and thermal responses can take multiple samples after dosing or cooling, whereas isolated spikes should not immediately change the trusted state.

The state update can be written as

$$z_s(t) = \mathcal{T}(z_s(t-1), g_s(t), q_s(t), h_s(t), a(t)), \quad (10)$$

Table 3: Manuscript AASVR-R configuration and exhaustive sensitivity ranges. The exhaustive profile is implemented as a shardable reproducibility protocol; the manuscript tables use the held-out benchmark unless explicitly marked as exhaustive sensitivity.

Parameter family	Manuscript setting	Exhaustive sensitivity range
Robust scale c_s	6.0	2, 3, 4, 5, 6, 8, 10, 12
Trust threshold $q_{\min}$	0.70	0.50, 0.60, 0.70, 0.80, 0.90
Transient / persistent limits	2 / 5 samples	transient 1, 2, 3, 5; persistent 3, 5, 8, 12
Authorization persistence k_s	3 samples	1, 2, 3, 5, 8
Replay cooldown	4 samples	0, 1, 2, 4, 8, 16, 40
Rectification branch	hold	hold, bounded prediction
CUSUM (α_c, h_c)	(0.10, 0.60)	$\alpha_c \in [0, 0.80]$, $h_c \in [0.20, 5.00]$
Response memory λ	0.50	0.50, 0.70, 0.80, 0.90, 0.95
Risk thresholds	0.75 / 0.65 / 0.50	high 0.60–0.90, medium 0.50–0.80, low 0.30–0.60
Response window	sensor-specific 6–8	4, 6, 8, 12, 20 samples

where $h_s(t)$ summarizes persistence-window evidence. This representation makes the state machine explicit without treating every transition as a separate hand-tuned rule in the manuscript. The transition function $\mathcal{T}$ is implemented by Algorithm 1.




Figure 4: AASVR-R persistence state machine.

5.5. Rectification and actuation authorization

When a reading is accepted, AASVR-R sets $\hat{x}_s(t) = y_s(t)$. When it is suspect, the method either holds $\hat{x}_s(t-1)$ or applies a bounded update

limited by $r_s \Delta t$:

$$\hat{x}_s(t) = \begin{cases} y_s(t), & z_s(t) \in \{\text{normal, accepted}\}, \\ \hat{x}_s(t-1), & z_s(t) = \text{suspect}, \\ \hat{x}_s(t-1) + \text{clip}(\Delta y_s, -r_s \Delta t, r_s \Delta t), & \text{bounded}, \\ \text{alert}, & z_s(t) = \text{fault}. \end{cases} \quad (11)$$

Actuation is authorized only if the trusted value is outside the control band, the violation persists for the configured number of samples, the sensor state is normal or accepted-regime-shift, the trust score is sufficient, the actuator is not in cooldown, and the recent response reliability $\eta_s(t)$ is above the risk-specific threshold:

$$A_s(t) = \mathbf{1}\{\hat{x}_s(t) \notin [L_s^a, U_s^a]\} \mathbf{1}\{p_s(t) \geq k_s\} \mathbf{1}\{q_s(t) \geq q_{\min}\} \\ \times \mathbf{1}\{\eta_s(t) \geq \eta_{\min, \rho(s)}\} \mathbf{1}\{z_s(t) \in \mathcal{Z}_{\text{safe}}\} \mathbf{1}\{\ell_s(t) = 0\}. \quad (12)$$

Here $p_s(t)$ is the number of consecutive trusted control-band violations, k_s is the persistence requirement, $\rho(s)$ is the action-risk class, $\mathcal{Z}_{\text{safe}} = \{\text{normal, accepted}\}$, and $\ell_s(t)$ is the actuator lockout indicator. For pH and EC, this rule complements the 10-minute dosing lockout already present in the deployed system; for CO2 and temperature, it prevents transient readings from becoming immediate gas-release or cooling commands.

5.6. Actuator-response residual

After an authorization at time t_0 , AASVR-R checks whether the trusted measurement moves in the expected direction within a response window. Let $d_s(t_0) \in \{-1, 0, +1\}$ denote the expected direction: high pH or EC implies a decreasing response after acid or dilution, low pH or EC implies an increasing response after base or nutrient dosing, and cooling implies decreasing temperature. If $d_s(t_0) = 0$, the response gate is not evaluated and no response-reliability update is made for that episode. Otherwise, the response evidence is

$$M_s(t_0) = \max_{\tau_{\min} \leq \tau \leq \tau_{\max}} d_s(t_0) [\hat{x}_s(t_0 + \tau) - \hat{x}_s(t_0)]. \quad (13)$$

The response gate is

$$G_s^{AR}(t_0) = \mathbf{1}\{M_s(t_0) \geq \rho_s\}. \quad (14)$$

If this gate fails, the method logs an actuator-response residual, withholds further authority from the sensor according to the reliability memory, and may escalate to fault alert. The reliability memory is

$$\eta_s(t+1) = \lambda\eta_s(t) + (1-\lambda)c_s(t), \quad (15)$$

where $c_s(t) = 1$ after a confirmed response and $c_s(t) = 0$ after a failed response. This term is intentionally different from $q_s(t)$: $q_s(t)$ scores the current sample, while $\eta_s(t)$ scores whether recent actuation episodes produced credible sensor response.

Table 4: Response-verification parameters used for the hydroponic primary sensors. Windows are sample counts at the 15–16 s replay cadence. Bidirectional sensors choose the sign from the trusted value relative to the control band.

Sensor	Actuator channel(s)	Window	Min. response	Risk
pH	acid/base doser	6	0.04 pH	High
EC	nutrient/water pump	6	30 $\mu\text{S cm}^{-1}$	High
Air temperature	fan/cooler	8	0.30 °C	Medium
Humidity	humidifier/dehumidifier	8	2.0% RH	Medium
Water temperature	water chiller	8	0.25 °C	Medium
CO2	CO2 valve/ventilation fan	6	80 ppm	Medium

Bounded-update observation. Assume the true process state satisfies $|x_s(t) - x_s(t-1)| \leq r_s\Delta t$ and the previous trusted estimate has error $|\hat{x}_s(t-1) - x_s(t-1)| \leq e_{t-1}$. Under the bounded rectification branch in Eq. (11), $|\hat{x}_s(t) - x_s(t)| \leq e_{t-1} + 2r_s\Delta t$ for any raw value $y_s(t)$, including arbitrarily large excursions. The result follows from applying a clipped update of magnitude at most $r_s\Delta t$ while the true state also moves by at most $r_s\Delta t$. Under a sustained suspect state, this per-sample bound can accumulate linearly, which is the formal counterpart of the empirical drift weakness. This is therefore a bounded-input statement for the controller-facing estimate, not a full closed-loop stability proof.

Algorithm 1 Streaming AASVR-R update for one sensor

Require: Raw value $y_s(t)$, previous trusted state $\hat{x}_s(t-1)$, recent history, limits, actuator state.

- 1: Estimate robust local scale $\sigma_{\Delta,s}(t)$ and tolerance $\xi_s(t)$.
 - 2: Apply physical-range, rate, trusted-deviation, missingness, trend, and stuck-value gates.
 - 3: If actuator response is pending, check the response-window residual and update $\eta_s(t)$.
 - 4: Update the persistence state from the gate outcome.
 - 5: **if** the reading is accepted **then**
 - 6: Set $\hat{x}_s(t) \leftarrow y_s(t)$.
 - 7: **else if** the reading is transient suspect **then**
 - 8: Hold or bounded-predict $\hat{x}_s(t)$.
 - 9: **else**
 - 10: Keep the last safe estimate and log a fault-alert decision.
 - 11: **end if**
 - 12: Authorize actuation only if trusted state, persistence, score, reliability, and cooldown rules are satisfied.
 - 13: **return** Trusted value, sensor state, trust score, actuation flag, and log record.
-

6. Evaluation Protocol

6.1. Hydroponic deployment data

The hydroponic data contain 121244 raw rows collected from 2024-02-27 to 2024-03-28 at a median cadence of approximately 16 seconds. The primary analysis uses EC, pH, humidity, air temperature, water temperature, and CO2. Water level is part of the deployed system but is excluded from the primary analysis because the available raw values are not defensibly calibrated.

6.2. Agronomic linkage protocol

Agronomic measurements are treated as deployment-context outcomes rather than as causal evidence for the algorithm. The per-plant harvest file contains 63 records across two experiments and three treatments. The analysis uses experiment, treatment, plant identifier, fresh and dry mass of aerial part, total plant length, root length, shoot length, total leaf count,

and leaves longer than 10 cm. Total leaf count and leaves longer than 10 cm are counts, not length measurements. The linkage analysis connects P1 harvest outcomes to measurement-to-actuation exposure metrics from the hydroponic replay: unsafe-band residence, untrusted-sample rate, alert rate, authorization rate, response-residual rate, and raw-to-trusted instability. P2 and P3 remain agronomic comparators only because they did not generate AASVR-R decision streams. This design supports the mechanistic claim that AASVR-R reduces unreliable control exposure during a complete growth cycle; it does not isolate AASVR-R from enclosure, lighting, cooling, CO₂ regulation, nutrient delivery, or automation.

6.3. Supplementary external stress tests

HAI 21.03 and SKAB are retained in the reproducibility outputs as supplementary stress tests rather than as main evidence. Their labels describe cyber-physical attacks and process-loop anomalies, not hydroponic sensor faults. The direct-transfer results are therefore reported only as a limitation check.

6.4. Fault injection protocol

Synthetic faults are injected into windows sampled from the real hydroponic background. For each sensor, fault type, magnitude, and duration, three deterministic replicates are generated and assigned to tune, validation, and test splits. The reported hydroponic benchmark uses only the held-out test split; AASVR-R and baseline hyperparameters are selected on validation trials. Fault windows include 90 pre-fault samples and 120 post-fault samples so that detection delay and false alarms are evaluated around the injected event rather than in isolation.

To address sensitivity to fault-design choices, the repository now implements a full factorial exhaustive profile with six primary sensors, fourteen fault classes, four robust-scale magnitude levels, eight durations, and 50 disjoint insertion windows per cell for each of tune, validation, and test splits. This produces 403200 protocol-defined trials before method expansion. The additional fault classes include plausible and implausible stuck values, high/low saturation, nonlinear drift, multiplicative gain drift, and intermittent bursts. The exhaustive profile is shardable and resumable; it is intended to replace the compact manuscript benchmark when the full run is executed for final resubmission, while the compact held-out benchmark remains the currently reported result set.

Table 5: Compact synthetic fault-injection protocol for the current hydroponic replay benchmark. The exhaustive reproducibility profile extends this table with plausible/implausible stuck values, high/low saturation, nonlinear drift, multiplicative gain drift, intermittent bursts, more durations, more magnitudes, and 50 disjoint insertion windows per cell.

Fault type	Injection rule over fault window $\mathcal{I}$	Purpose
Spike	$y'_s(t) = y_s(t) + m$	Single-window offset
Multi-spike	$y'_s(t) = y_s(t) + b_t m, b_t \in \{-1, 1\}$	Alternating impulse faults
Dropout	$y'_s(t) = \text{NaN}$	Missing samples
Drift	$y'_s(t) = y_s(t) + m(t - t_0)/ \mathcal{I} $	Slow monotonic deviation
Bias	$y'_s(t) = y_s(t) + m$	Persistent offset
Noise burst	$y'_s(t) = y_s(t) + \epsilon_t, \epsilon_t \sim \mathcal{N}(0, m^2)$	Increased variance
Step	$y'_s(t) = y_s(t) + m$	Abrupt sustained change
Stuck-at	$y'_s(t) = y_s(t_0)$	Constant repeated value
Saturation	$y'_s(t) = m$	Sensor rail or implausible fixed value

6.5. Counterfactual response replay

A second replay evaluates the actuator-response residual. Replay-derived controller authorizations are treated as pseudo-actuator commands because the available feed does not contain real actuator-state columns. Healthy trials inject a first-order response in the expected direction after authorization. Response-fault trials inject no-response, stuck-after-command, weak-response, delayed-response, and wrong-direction behavior. These trials test whether the supervisor detects missing measurement response and suppresses repeated authority; they are not interpreted as measured physical closed-loop disturbances.

6.6. Prospective limitation-closing evidence

The reproducibility package now treats three missing evidence sources as explicit optional schemas rather than informal future work. First, a real actuator-state log should record timestamped actuator identity, commanded and measured state, command source, target sensor, expected response direction, dose or runtime, lockout status, and manual override status. This converts the response residual from counterfactual replay into measured command-response evidence. Second, an independent reference-measurement file should pair raw sensor values with a calibrated handheld or benchtop reference instrument, including device identity and operator notes. This supports direct scoring of plausible bias, drift, and stuck-at readings. Third, water-level reintegration requires paired raw ultrasonic readings and

physical reference levels in centimeters. The current reproducibility utility fits a linear calibration, $\ell_{\text{cm}} = ay_{\text{raw}} + b$, and reports slope, intercept, RMSE, and R^2 ; water level should remain excluded until those diagnostics are acceptable for the physical tank geometry. Missing optional files are allowed in the current pipeline and produce absence-status tables; they do not change the held-out hydroponic benchmark.

6.7. Baselines and metrics

The baselines are raw thresholding, moving average, moving median, Hampel, Kalman, EWMA, CUSUM, PCA-style monitoring, GLR, recursive PCA, Isolation Forest, One-Class SVM, and local outlier factor. A duplicated physical-threshold rule remains in the reproducibility outputs but is not interpreted as a separate scientific baseline. Every displayed baseline is tuned on the validation split using Eq. (2) and then wrapped in the same persistence/cooldown authorization supervisor for false-actuation scoring. The hydroponic synthetic-fault benchmark reports precision, recall, specificity, balanced accuracy, detection delay, false alarm events, false actuation mean, and false actuation standard deviation.

Bootstrap confidence intervals use 1000 resamples of injected-fault trials, not individual samples. Paired Wilcoxon signed-rank tests compare AASVR-R against each baseline on per-trial balanced accuracy and false actuation, with Holm correction applied separately across the 13-baseline balanced-accuracy family and the 13-baseline false-actuation family. Rank-biserial correlation and paired d_z are reported in the reproducibility tables, and a sign test is also reported for false-actuation differences because that distribution is zero-inflated and skewed. Sensor-fault grids use Benjamini-Hochberg correction.

Table 6: Baseline taxonomy used in the hydroponic benchmark. The final column marks whether a method explicitly separates anomaly detection from actuation authorization.

Family	Methods	Primary role	Act.-aware
Rule	Raw threshold	Existing physical-limit control logic	No
Filters	Moving average, moving median, Hampel, Kalman	Noise and spike suppression	No
Process monitoring	EWMA, CUSUM, PCA, GLR, recursive PCA	Shift or process-deviation detection	No
One-class ML	Isolation Forest, One-Class SVM, LOF	Data-driven anomaly boundary	No
Proposed method	AASVR-R	Validation, rectification, authorization, and response verification	Yes

7. Results

The results are organized as follows. Section 7.1 characterizes the hydroponic streams. Section 7.2 reports the held-out synthetic-fault benchmark and paired tests. Section 7.3 breaks performance down by sensor and fault class. Section 7.4 separates the contribution of gating, rectification, and authorization. Section 7.5 reports ablations and tuning sensitivity. Section 7.6 links per-plant agronomic outcomes to replay-derived exposure metrics, and Section 7.7 connects the aggregate benchmark to representative deployment traces.

7.1. Hydroponic deployment characterization

The deployment data are first used to characterize the real measurement background on which AASVR-R operates. The 121244-row feed covers a continuous 30-day lettuce growth period, with approximately 16-second median sampling. The key point for this paper is not that the dashboard collected data, but that the closed-loop system generated the kind of sensor-to-actuator stream in which a bad reading can trigger a physical correction.

7.2. Hydroponic synthetic-fault benchmark

Table 8 summarizes the primary target-domain result on 270 held-out synthetic-fault trials. With the validation-selected Page/CUSUM residual

Table 7: Hydroponic measurement streams used in the primary AASVR-R analysis. The table separates deployed sensing from quantitative use so that the water-level exclusion is explicit rather than hidden.

Variable	Module	Primary analysis	Control relevance
EC	Solution	Yes	Nutrient/water dosing evidence
pH	Solution	Yes	Acid/base dosing evidence
Humidity	Growth	Yes	Growth-environment monitoring
Air temperature	Growth	Yes	Air-cooling evidence
Water temperature	Solution	Yes	Chiller evidence
CO2	Growth	Yes	CO2 valve evidence
Water level	Solution	No	Warning channel; calibration unresolved

gate, AASVR-R is close to the base AASVR detector on balanced accuracy and exceeds the displayed non-AASVR baselines while retaining low false-authorized actuation. LOF has higher recall, but it pays for this with lower specificity and more than three times as many alert samples. The base AASVR detector is fractionally higher in balanced accuracy than AASVR-R because it does not withhold repeated authority after failed response, but its false-authorized actuation count is substantially larger. One-Class SVM also has low false-authorized actuation, but it obtains this by rejecting aggressively, with a false-positive rate of 0.667 and nearly 890 alert samples per trial. The contribution claim is therefore a combined one: the standard CUSUM residual gate strengthens detection, while AASVR-R remains structurally auditable as a controller-facing reliability layer.

Table 8: Held-out hydroponic synthetic-fault benchmark. Baselines are validation-tuned and evaluated on the same test split with a shared persistence/cooldown supervisor. Bold entries mark the best value in each metric column among the displayed methods; lower is better for false actuation.

Method	Recall	Specificity	Bal. acc.	FPR	Delay	False act.	False act. SD
AASVR	0.821	0.870	0.845	0.130	0.000	0.196	0.738
AASVR-R	0.821	0.867	0.844	0.133	0.000	0.041	0.290
LOF	0.853	0.798	0.826	0.202	0.270	0.130	0.629
Kalman	0.613	0.803	0.708	0.197	0.000	0.515	1.175
Isolation Forest	0.581	0.826	0.703	0.174	0.126	0.630	1.277
Hampel	0.411	0.984	0.698	0.016	0.000	1.041	1.502
Moving median	0.411	0.984	0.698	0.016	0.000	0.993	1.496
Recursive PCA	0.317	0.988	0.652	0.012	0.000	1.319	1.653
One-Class SVM	0.912	0.333	0.623	0.667	0.089	0.070	0.403
GLR	0.129	0.999	0.564	0.001	0.000	1.330	1.638

Table 9: Paired held-out effects for AASVR-R against representative methods. Positive Δ BA means AASVR-R has higher per-trial balanced accuracy; positive Δ false actuation means the comparison method has more false actuations than AASVR-R. d_z is the paired standardized BA effect size; $p_{\text{FA,sign}}$ is Holm-corrected for the sign-test false-actuation family.

Baseline	Δ BA	d_z	p_{BA}	Δ False act.	95% CI	$p_{\text{FA,sign}}$
Base AASVR	-0.001	-0.719	1.000	0.156	[0.081, 0.241]	< 0.001
LOF	0.018	0.113	< 0.001	0.089	[0.030, 0.152]	0.021
One-Class SVM	0.221	1.240	< 0.001	0.030	[-0.026, 0.081]	0.212
Kalman	0.136	0.594	< 0.001	0.474	[0.341, 0.607]	< 0.001
Hampel	0.147	0.646	< 0.001	1.000	[0.822, 1.185]	< 0.001
CUSUM	0.204	0.946	< 0.001	1.233	[1.044, 1.426]	< 0.001

The base-AASVR balanced-accuracy effect size in Table 9 is large despite a small mean difference because the per-trial difference is highly consistent (standard deviation about 0.0018). It is therefore not interpreted as an operationally large detection advantage; the operationally relevant contrast is the false-authorized actuation reduction.

Figure 5 and Table 10 show the audit trail for a held-out pH fault replay. The raw pH value jumps from the trusted state near 6.09 to a faulted value above 8.0. The physical and rate gates remain true, but the trusted-deviation gate fails. The supervisor holds the trusted value, lowers the trust score, transitions from suspect transient to persistent deviation and then fault alert, and never authorizes actuation from the untrusted pH evidence.

Table 10: Audit-trail excerpt for the pH replay in Fig. 5. G^D is the trusted-deviation gate.

Rel. sample	Raw	Trusted	G^D	State	Action	Authorized
-1	6.09	6.09	True	normal	accept	False
0	8.09	6.09	False	suspect	hold	False
2	8.09	6.09	False	persistent	hold	False
4	8.09	6.09	False	fault	alert	False
8	8.09	6.09	False	fault	alert	False




Figure 5: AASVR-R audit trail for a held-out pH fault replay. The shaded region is the injected fault window; rejected samples are logged without authorizing actuation.




Figure 6: Held-out hydroponic synthetic-fault balanced accuracy across methods. AASVR-R combines robust gates with a standard CUSUM residual test and response-aware authorization, giving a strong balanced-accuracy/false-actuation tradeoff.




Figure 7: Bootstrap uncertainty for balanced accuracy. Intervals are trial-level resamples.

7.3. Sensor-wise and fault-type comparisons

Table 11 breaks down AASVR-R performance by hydroponic sensor. The CUSUM residual gate substantially improves the EC stream relative to the earlier single-window gate, so all primary sensors now remain above 0.83 balanced accuracy. CO2 remains the strongest stream. EC and pH are still the most operationally sensitive because dosing actions can change the nutrient solution, but the response-memory layer keeps their false-authorized actuation low in replay.

Table 12 and Figure 8 show performance by injected fault class. AASVR-R performs best on spike, saturation, dropout, step, bias, noise-burst, and multi-spike faults. The sequential CUSUM residual gate improves slow-shift and EC-bias behavior but does not solve frozen plausible values. Stuck-at remains the weakest detection case because an exactly repeated value can be indistinguishable from a genuinely steady process unless there is independent excitation, reference instrumentation, or a command-response episode.

Table 11: AASVR-R synthetic-fault performance by hydroponic sensor. False actuation is reported as mean events per injected-fault trial.

Sensor	Recall	Specificity	Bal. acc.	False act.
CO2	0.854	0.878	0.866	0.056
Water temperature	0.821	0.870	0.846	0.037
Air temperature	0.823	0.860	0.841	0.093
pH	0.808	0.863	0.836	0.019
EC	0.801	0.863	0.832	0.000

Table 12: AASVR-R synthetic-fault performance by injected fault type. Stuck-at and drift are the most difficult conditions because frozen plausible values and slow monotonic movement can resemble real process behavior.

Fault type	Recall	Specificity	Bal. acc.	False act.
Spike	0.996	0.875	0.935	0.000
Saturation	1.000	0.862	0.931	0.000
Dropout	1.000	0.854	0.927	0.000
Step	0.976	0.875	0.925	0.000
Bias	0.983	0.866	0.925	0.000
Noise burst	0.936	0.873	0.904	0.033
Multi-spike	0.932	0.867	0.900	0.000
Drift	0.496	0.861	0.678	0.067
Stuck-at	0.074	0.869	0.472	0.267




Figure 8: Hydroponic synthetic-fault heatmap by sensor and fault type. The CUSUM residual gate improves persistent shifts, while stuck-at behavior remains the clearest limitation.

7.4. Component contribution and ablation

Table 13 separates the contribution of the supervisory components on the same 270 held-out test trials as Table 8. Raw thresholding with the shared supervisor has high specificity but poor detection because only physically extreme values are rejected. Adding robust gates and hold rectification raises balanced accuracy and reduces false-authorized actuation. The trust-score-only row is identical to robust gates plus hold, showing that in the selected configuration the trust score is mainly diagnostic unless it is coupled to the full persistence and authorization logic. The CUSUM residual gate is the main reason the final detector now surpasses LOF on balanced accuracy; the response-reliability layer is the main reason AASVR-R keeps false-authorized actuation low.

Table 13: Component-contribution analysis on the full 270-trial held-out synthetic-fault test split. Values are mean events or scores per trial.

Variant	Recall	Specificity	Bal. acc.	False act.
Raw threshold + supervisor	0.129	0.999	0.564	1.330
Robust gates + hold	0.737	0.887	0.812	0.363
Robust gates + trust + hold	0.737	0.887	0.812	0.363
Full AASVR	0.821	0.870	0.845	0.196
Full AASVR-R	0.821	0.867	0.844	0.041
AASVR without cooldown	0.821	0.870	0.845	0.411
AASVR without confirmation/cooldown	0.821	0.869	0.845	1.730

Table 14 shows the internal AASVR-R ablation results on a deterministic 120-trial subset of the held-out test split. It is not numerically identical to the full 270-trial reference in Tables 8 and 13, and the caption states the subset explicitly to avoid interpreting small differences as algorithm changes. Several ablations preserve similar detection scores but change control behavior. Removing actuation confirmation increases false actuation from 0.017 to 0.300 events per trial in that subset, removing cooldown increases it to 0.233, and disabling response reliability increases it to 0.158. The stuck-detector contrast is better interpreted through the stuck-at-specific result in Table 12; its aggregate effect is small because stuck-at is one of several injected fault classes.

The validation sweep in Table 15 shows that the selected AASVR-R configuration is not an isolated optimum. Multiple trust-score and transient-limit settings produced the same validation objective when the robust scale

Table 14: AASVR-R ablation summary on sampled held-out test trials. Rows are ordered to emphasize false-actuation impact relative to the full method. Detection scores can remain similar while control outcomes change, which supports evaluating control-aware metrics rather than point-wise detection alone.

Variant	Bal. acc.	Recall	False act.
No actuation confirmation	0.845	0.818	0.300
No cooldown	0.845	0.818	0.233
No response reliability	0.845	0.818	0.158
No trend guard	0.828	0.783	0.033
Full AASVR-R	0.843	0.818	0.017
No stuck detector	0.845	0.811	0.067
No robust MAD scale	0.843	0.847	0.008
Bounded prediction mode	0.789	0.703	0.050

multiplier was 6.0. These parameters were selected before evaluating the held-out test split in Table 8. A transfer check over the top validation settings showed held-out balanced accuracy within 0.0011 across the top 15 settings on a fixed sampled test subset; objective-weight profiles that increased the false-actuation penalty selected the same top setting in the current grid. This indicates that the reported result is not sensitive to the arbitrary choice among tied $q_{\min}$ or transient-limit settings.

Table 15: Top AASVR-R parameter settings from the validation split. The selected setting uses $c_s = 6.0$, $q_{\min} = 0.70$, and transient limit 2; tied settings are retained to show sensitivity.

c_s	$q_{\min}$	Transient limit	Bal. acc.	Objective
6.0	0.70	2	0.824	0.744
6.0	0.70	1	0.824	0.744
6.0	0.55	3	0.824	0.744
6.0	0.55	2	0.824	0.744
6.0	0.55	1	0.824	0.744

7.5. Agronomic deployment linkage

Table 16 reports the per-plant harvest summary from the two lettuce experiments. P1 has the largest fresh aerial mass and NL10 count in both

experiments, but this comparison is not interpreted causally because P1 differs from P2 and P3 in enclosure, lighting, cooling, CO2 regulation, nutrient delivery, and automation. The large P2/P3 reduction in Experiment 2 is treated as platform/treatment confounding or separate growth stress, not as an algorithmic effect. Instead, the agronomic data are used to verify that the measurement-to-actuation pipeline operated through a complete crop cycle and to attach plant-level context to replay-derived control-risk exposure.

Table 16: Agronomic deployment context from the per-plant harvest file. Values are mean with bootstrap 95% confidence intervals. TNL and NL10 are counts.

Experiment	Treatment	FMAP (g)	TPL (cm)	NL10 (count)
1	P1	252.87 [236.74, 270.13]	81.01 [79.85, 82.02]	10.21 [8.57, 12.00]
1	P2	141.98 [133.22, 150.98]	57.27 [49.80, 64.06]	9.50 [7.71, 11.29]
1	P3	94.09 [88.75, 99.42]	20.67 [19.67, 21.67]	7.07 [6.00, 8.29]
2	P1	264.81 [243.49, 291.51]	72.52 [64.92, 80.14]	13.00 [11.71, 14.43]
2	P2	68.24 [54.94, 81.33]	19.46 [17.19, 21.90]	3.00 [2.29, 3.86]
2	P3	109.49 [95.33, 126.83]	20.12 [18.77, 21.59]	7.57 [6.57, 8.57]

The exposure linkage in Table 17 is restricted to P1 because only P1 produced the sensor-to-actuator decision stream used by AASVR-R. The table now reports authorization counts and missed authorization opportunity rather than treating a low replay authorization rate as automatically beneficial. On the full deployment replay, AASVR-R authorizes only 11 sensor-sample decisions while the trusted state remains outside the control band for much of the replay. This means the full-feed replay should not be interpreted as measured dosing reduction or as evidence that the deployed crop cycle required almost no actuation; it instead exposes a limitation of scoring response-aware authorization without real actuator-state logs and a deployment-matched command history.

Table 17: P1 agronomic linkage to replay-derived measurement-to-actuation exposure. Rates are fractions of sensor-sample decisions over the hydroponic feed; low authorization without measured actuator logs is interpreted as replay silence, not as operational superiority.

Method	Untrusted	Alert	Auth. count	Auth. rate	Missed opp.	Resp. residual
Raw threshold	0.001765	0.001765	192506	0.264626	0.000000	0.000000
AASVR	0.217791	0.147743	60723	0.083472	0.704638	0.000000
AASVR-R	0.217796	0.147743	11	0.000015	0.788095	0.000005

7.6. Counterfactual response replay

Table 18 reports the actuator-response replay as a capability demonstration for AASVR-R, not as an independent comparative benchmark. The response faults are deliberately defined relative to the method’s expected-response model, so the result should be read as evidence that the response-residual machinery is wired correctly and as a guide to which response classes it can expose. AASVR-R detects no-response and stuck-after-command cases consistently, detects some delayed and weak responses, and preserves all healthy-response trials in this replay. Wrong-direction faults are usually suppressed by existing plausibility and reliability logic rather than explicitly labeled as response residuals, so they remain a target for stronger direction-specific deployment logs.

Table 18: AASVR-R counterfactual actuator-response capability replay. Response correctness is one for healthy trials without false residuals and for response-fault trials with a detected residual. Authorizations are replay-derived controller commands, not measured physical actuator events.

Response class	Fault detection	Correctness	Delay	Authorizations	Final η_s
Healthy	–	1.000	–	1.000	1.000
No response	1.000	1.000	11.600	1.400	0.500
Stuck after command	1.000	1.000	11.600	1.400	0.500
Weak response	0.600	0.600	19.000	3.200	0.700
Delayed response	0.600	0.600	8.000	1.800	0.700
Wrong direction	0.000	0.000	–	1.000	0.500

Table 19 reports a second counterfactual replay in which each sensor receives a safe diagnostic pulse while its value is inside the normal control band. Healthy trials inject a small response in the expected direction; stuck-during-diagnostic trials keep the measurement frozen; weak-response trials move less than the response threshold. This experiment is not claimed as evidence from the existing deployment logs. It quantifies the value of a future actuator-state channel: once a diagnostic action is known, a plausible frozen value becomes detectable through missing response rather than through passive range/rate gates alone.

7.7. Representative hydroponic trace

Figure 10 shows replay segments from the hydroponic deployment. The traces are included to connect the aggregate benchmark to the actual deploy-

Table 19: Diagnostic-excitation replay for future deployments with actuator-state logging. Correctness is one for healthy trials without false residuals and for response-fault trials with a detected residual.

Method	Fault detection	Correctness	Delay	Final η_s
AASVR-R	1.000	1.000	12.000	0.667




Figure 9: Counterfactual pH no-response replay. The scenario is initialized above the pH control band so that AASVR-R can authorize correction, open a response window, and flag a response residual when the trusted pH fails to decrease.

ment setting. They illustrate the measurement-to-state transformation that precedes any actuation decision.




Figure 10: Representative hydroponic replay traces. The raw streams and trusted estimates are plotted over early deployment segments. These traces are not used as ground-truth fault labels; they show how the method exposes a trusted state rather than passing every raw measurement directly to the controller.

8. Discussion

8.1. Why actuation awareness matters

The ablation, component, response-replay, diagnostic-excitation, and audit-trail results show why actuation-aware validation should be evaluated with control metrics and decision transparency. After adding the validation-selected Page/CUSUM residual gate, AASVR-R exceeds LOF on balanced accuracy while keeping substantially fewer alert samples and lower false-authorized actuation. The detector gain comes from a standard sequential residual test rather than from a learned one-class boundary, so the novelty is not the CUSUM detector itself but its placement inside an auditable measurement-to-actuation supervisor. Each rejection remains traceable to a range/rate/deviation, trend, CUSUM, stuck-value, or response-residual condition. The trust-score-only component result is also informative. In this configuration it does not change the robust-gate baseline because the binary gates already determine most accept/reject decisions; its value appears mainly as a diagnostic and as an input to the full state/authorization supervisor. For hydroponic systems, these distinctions are practical because a logged anomaly, a held trusted state, a pump command, a diagnostic pulse, and a failed post-command response are different engineering decisions.

8.2. *External transfer*

The supplementary HAI and SKAB checks show that the same streaming interface can run on independent process datasets, but direct parameter transfer is weak. HAI attacks and SKAB process anomalies are not identical to hydroponic sensor faults, and the current method uses hydroponic physical ranges and rate assumptions. These results are therefore not used as evidence of cross-domain dominance; they identify the need for process-specific limits or model residuals before claiming industrial transfer.

8.3. *Limitations*

The hydroponic experiment does not provide independent reference instruments for every sensor at every time point. Synthetic fault injection is therefore necessary for controlled scoring, but injected faults cannot cover every real failure mode. The reported false-actuation values are replayed authorization events, not measured physical process disturbances, because the replay benchmark does not close the loop around real actuators. The response and diagnostic-excitation replays are also counterfactual: replay-derived authorizations and safe diagnostic pulses stand in for actuator commands because the available feed does not contain real actuator-state columns.

The remaining limitations are therefore identifiable and measurable. Stuck-at behavior and weak response remain the principal algorithmic weaknesses because they can resemble real steady nutrient-solution dynamics without reference instrumentation or measured actuator effects. AASVR-R reduces repeated authority after failed response and improves slow-shift detection through the CUSUM residual gate, and the diagnostic replay shows how future actuator-state logging can expose plausible frozen values through missing response evidence. It still does not make frozen plausible values fully detectable when no command-response, diagnostic episode, or independent reference sample occurs. Water level is excluded from primary analysis because the available raw values are not defensibly calibrated; the repository now includes a calibration schema and fitting script that should be satisfied before reintegration. Crop outcomes can be linked to replay-derived exposure metrics as full-cycle deployment context, but they remain non-causal because enclosure, lighting, cooling, nutrient delivery, carbon dioxide regulation, and automation are confounded.

9. Conclusion

We presented a closed-loop hydroponic control system and proposed AASVR-R as its response-aware sensor validation and rectification layer. The method separates validation, state update, rectification, actuation authorization, and response verification so that untrusted measurements can be logged without automatically driving physical commands. On held-out hydroponic synthetic faults, the standard CUSUM residual gate makes AASVR-R competitive with tuned detectors, while the response-aware supervisor retains low false-authorized actuation and a transparent decision structure. Counterfactual response and diagnostic-excitation replay serve as capability demonstrations showing that actuator-response residuals can expose no-response, weak-response, and stuck-after-command cases that single-stream gates miss. The remaining requirement is not another passive filter alone, but real actuator-state logs and reference measurements that make command-response evidence observable during future deployments.

Data Availability

The repository contains code, configuration files, processed-data descriptors, Docker workflow files, and paper-generation assets. Raw hydroponic and external benchmark files are not committed to the repository. Public datasets can be obtained from their original providers, and request-access datasets should be obtained under the providers' terms.

Declaration of Competing Interest

The authors declare no competing interests for the current manuscript draft.

Declaration of Generative AI and AI-Assisted Technologies

AI-assisted tools were used to help organize code, draft text, and check consistency. The authors remain responsible for the scientific content, claims, results, and final manuscript.

References

- [1] H. Y. Teh, A. W. Kempa-Liehr, K. I.-K. Wang, Sensor data quality: a systematic review, *Journal of Big Data* 7 (11) (2020). doi:10.1186/s40537-020-0285-1.
- [2] X. Zou, W. Liu, Z. Huo, S. Wang, Z. Chen, C. Xin, Y. Bai, Z. Liang, Y. Gong, Y. Qian, L. Shu, Current status and prospects of research on sensor fault diagnosis of agricultural internet of things, *Sensors* 23 (5) (2023) 2528. doi:10.3390/s23052528.
- [3] Y.-B. Lin, Y.-W. Lin, J.-Y. Lin, H.-N. Hung, Sensortalk: An iot device failure detection and calibration mechanism for smart farming, *Sensors* 19 (21) (2019) 4788. doi:10.3390/s19214788.
- [4] F. R. Hampel, The influence curve and its role in robust estimation, *Journal of the American Statistical Association* 69 (346) (1974) 383–393. doi:10.1080/01621459.1974.10482962.
- [5] R. E. Kalman, A new approach to linear filtering and prediction problems, *Journal of Basic Engineering* 82 (1) (1960) 35–45. doi:10.1115/1.3662552.
- [6] S. W. Roberts, Control chart tests based on geometric moving averages, *Technometrics* 1 (3) (1959) 239–250. doi:10.1080/00401706.1959.10489860.
- [7] E. S. Page, Continuous inspection schemes, *Biometrika* 41 (1/2) (1954) 100–115. doi:10.1093/biomet/41.1-2.100.
- [8] M. Basseville, I. V. Nikiforov, *Detection of Abrupt Changes: Theory and Application*, Prentice Hall, 1993.
- [9] R. J. Patton, P. M. Frank, R. N. Clark, *Fault Diagnosis in Dynamic Systems: Theory and Application*, Prentice Hall, 1989.
- [10] R. Isermann, *Fault-diagnosis systems: An introduction from fault detection to fault tolerance*, Springer, 2006. doi:10.1007/3-540-30368-5.

- [11] I. Hwang, S. Kim, Y. Kim, C. E. Seah, A survey of fault detection, isolation, and reconfiguration methods, *IEEE Transactions on Control Systems Technology* 18 (3) (2010) 636–653. doi:10.1109/TCST.2009.2026285.
- [12] H. Hotelling, Multivariate quality control, *Techniques of Statistical Analysis* (1947) 111–184.
- [13] J. E. Jackson, G. S. Mudholkar, Control procedures for residuals associated with principal component analysis, *Technometrics* 21 (3) (1979) 341–349. doi:10.1080/00401706.1979.10489779.
- [14] W. Li, H. H. Yue, S. Valle-Cervantes, S. J. Qin, Recursive pca for adaptive process monitoring, *Journal of Process Control* 10 (5) (2000) 471–486. doi:10.1016/S0959-1524(00)00022-6.
- [15] F. T. Liu, K. M. Ting, Z.-H. Zhou, Isolation forest, in: 2008 Eighth IEEE International Conference on Data Mining, IEEE, 2008, pp. 413–422. doi:10.1109/ICDM.2008.17.
- [16] B. Schölkopf, J. C. Platt, J. Shawe-Taylor, A. J. Smola, R. C. Williamson, Estimating the support of a high-dimensional distribution, *Neural Computation* 13 (7) (2001) 1443–1471. doi:10.1162/089976601750264965.
- [17] M. M. Breunig, H.-P. Kriegel, R. T. Ng, J. Sander, LOF: Identifying density-based local outliers, in: *Proceedings of the 2000 ACM SIGMOD International Conference on Management of Data*, 2000, pp. 93–104. doi:10.1145/342009.335388.
- [18] K. Tatas, A. Al-Zoubi, N. Christofides, C. Zannettis, M. Chrysostomou, S. Panteli, A. Antoniou, Reliable iot-based monitoring and control of hydroponic systems, *Technologies* 10 (1) (2022) 26. doi:10.3390/technologies10010026.
- [19] D. S. Domingues, H. W. Takahashi, C. A. Camara, S. L. Nixdorf, Automated system developed to control ph and concentration of nutrient solution evaluated in hydroponic lettuce production, *Computers and Electronics in Agriculture* 84 (2012) 53–61. doi:10.1016/j.compag.2012.02.006.

- [20] H.-K. Shin, W. Lee, J.-H. Yun, H. Kim, Hai 2021: A new dataset for anomaly detection in industrial control systems, in: Proceedings of the 14th Cyber Security Experimentation and Test Workshop, 2021, pp. 1–9. doi:10.1145/3474718.3474719.
URL <https://doi.org/10.1145/3474718.3474719>
- [21] I. D. Katser, V. O. Kozitsin, Skoltech anomaly benchmark (skab), <https://www.kaggle.com/dsv/1693952> (2020). doi:10.34740/KAGGLE/DSV/1693952.
- [22] Secure water treatment (swat) dataset, <https://www.sutd.edu.sg/itrust/itrust-labs/datasets/dataset-characteristics/swat/>, accessed 2026-05-18 (2026).
- [23] Water distribution (wadi) dataset, <https://www.sutd.edu.sg/itrust/itrust-labs/datasets/dataset-characteristics/wadi/>, accessed 2026-05-18 (2026).
- [24] Tennessee eastman reference data for fault-detection and decision support systems, https://data.dtu.dk/articles/dataset/Tennessee_Eastman_Reference_Data_for_Fault-Detection_and_Decision_Support_Systems/13385936, accessed 2026-05-18 (2026).
- [25] G. Niu, J. Masabni, Hydroponics, in: Plant Factory Basics, Applications and Advances, Academic Press, 2022, pp. 153–166. doi:10.1016/B978-0-323-85152-7.00014-6.
- [26] M. Asaduzzaman, G. Niu, T. Asao, Nutrients recycling in hydroponics: Opportunities and challenges toward sustainable crop production under controlled environment agriculture, *Frontiers in Plant Science* 13 (2022) 845472. doi:10.3389/fpls.2022.845472.